


Reliability: development (M06)



● Matched CNN (D0-M) □ Source-selected CNN △ Combined-loss CNN (D3)

Each panel pools context appearances (M01 spectra or M06 instrument-balanced masters) into up to ten equal-mass confidence bins.

Repeated appearances are not independent samples; no confidence interval or significance test is shown.

The pooled reliability curve is descriptive and differs from the registered mean per-context ECE; no extra calibration or temperature is fitted.

M01 scores individual spectra after three-seed probability averaging.

M06 averages probabilities within each instrument, then equally across instruments for a sample; raw spectra are not averaged.